



CSCI-GA.3033-004

Graphics Processing Units (GPUs): Architecture and Programming

Lecture : Parallel Patterns

Most slides of this
lecture are from:

- David Kirk/NVIDIA
and
- Wen-mei W. Hwu
EUniversity of
Illinois

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>




Convolution


Convolution

- An Array operation
- Output data element = **weighted sum** of a collection of neighboring input elements.
- The weights are defined by an input **mask array**.
- Usually used as filters to transform signals (or pixels or ...) into more desirable form.

Convolution



Convolution



Convolution can also be 2D.

Convolution

N


1	2	3	4	5	6	7
2	3	4	5	6	7	8
3	4	5	6	7	8	9
4	5	6	7	8	5	6
5	6	7	8	5	6	7
6	7	8	9	0	1	2
7	8	9	0	1	2	3

P

		321				

M

1	2	3	2	1
2	3	4	3	2
3	4	5	4	3
2	3	4	3	2
1	2	3	2	1



1	4	9	8	5
4	9	16	15	12
9	16	25	24	21
8	15	24	21	16
5	12	21	16	5

Convolution

```
__global__ void convolution_1D_basic_kernel(float *N, float *M, float *P,
int Mask_Width, int Width) {

    int i = blockIdx.x*blockDim.x + threadIdx.x;

    float Pvalue = 0;
    int N_start_point = i - (Mask_Width/2);
    for (int j = 0; j < Mask_Width; j++) {
        if (N_start_point + j >= 0 && N_start_point + j < Width) {
            Pvalue += N[N_start_point + j]*M[j];
        }
    }
    P[i] = Pvalue;
}
```

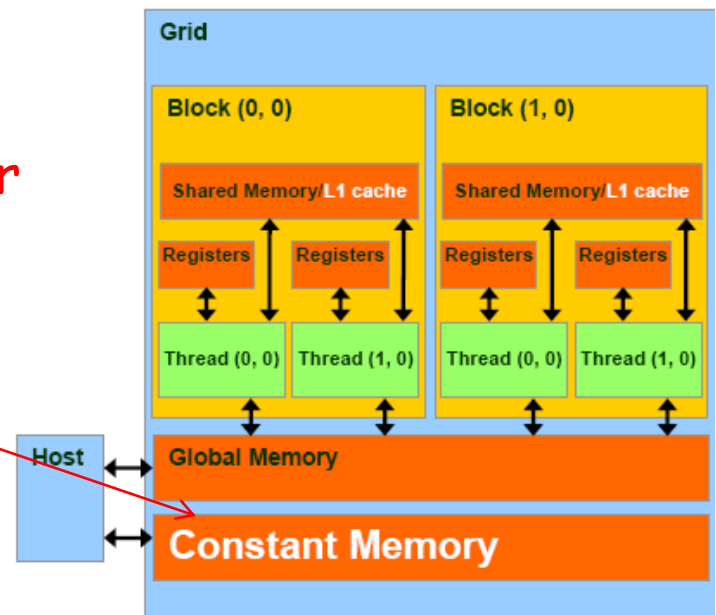
The 1D Version

- Thread organized as 1D grid.
- Pvalue allows intermediate values to be accumulated in registers to save DRAM bw.
- We assume *ghost* values are 0.
- There will be control flow divergence (due to ghost elements).
- **Ratio of floating point arithmetic calculation to global memory access is $\sim 1.0 \rightarrow$ What can we do??**

Regarding Mask M

- Size of M is typically small.
- The contents of M do not change during execution.
- All threads need to access M and in the same order.

Doesn't this make M a good candidate for constant memory?



Constant Memory

- Constant memory variables are visible to all thread blocks.
- Constant memory variables cannot be changed during kernel execution.
- The size of constant memory can vary from device to device.

How to Use Constant Memory

- Host code allocates, initializes variables the same way as any other variables that need to be copied to the device
- Use `cudaMemcpyToSymbol(dest, src, size)` to copy the variable into the device memory
- This copy function tells the device that the variable will not be modified by the kernel and can be safely cached.

Mask M and Constant Memory

- In host:
 - `#define MASK_WIDTH 10`
 - `__constant__ float M[MASK_WIDTH]`
 - Allocate and initialize a mask `h_M`
 - `cudaMemcpyToSymbol(M, h_M, MASK_WIDTH * sizeof(float),
offset, kind);`
- Kernel functions
 - access constant memory variables as global variables → no need to pass pointers of these variables to the kernel as parameter.

Question: Isn't the constant memory also in DRAM? Why is it assumed faster than global memory?

Answer:

- CUDA runtime knows that constant memory variables are not modified.
- It directs the hardware to aggressively **cache** them during kernel execution.

Reduction Trees

What? And Why?

- Arguably the most widely used parallel computation pattern.
- A commonly used strategy for processing large input data sets
 - There is no required order of processing elements in a data set (associative and commutative)
 - Partition the data set into smaller chunks
 - Have each thread to process a chunk
 - Use a reduction tree to summarize the results from each chunk into the final answer
- Google and Hadoop MapReduce frameworks are examples of this pattern

Reduction enables other techniques

- Reduction is also needed to clean up after some commonly used parallelizing transformations
- Example: Privatization
 - Multiple threads write into an output location
 - Replicate the output location so that each thread has a private output location
 - Use a reduction tree to combine the values of private locations into the original output location

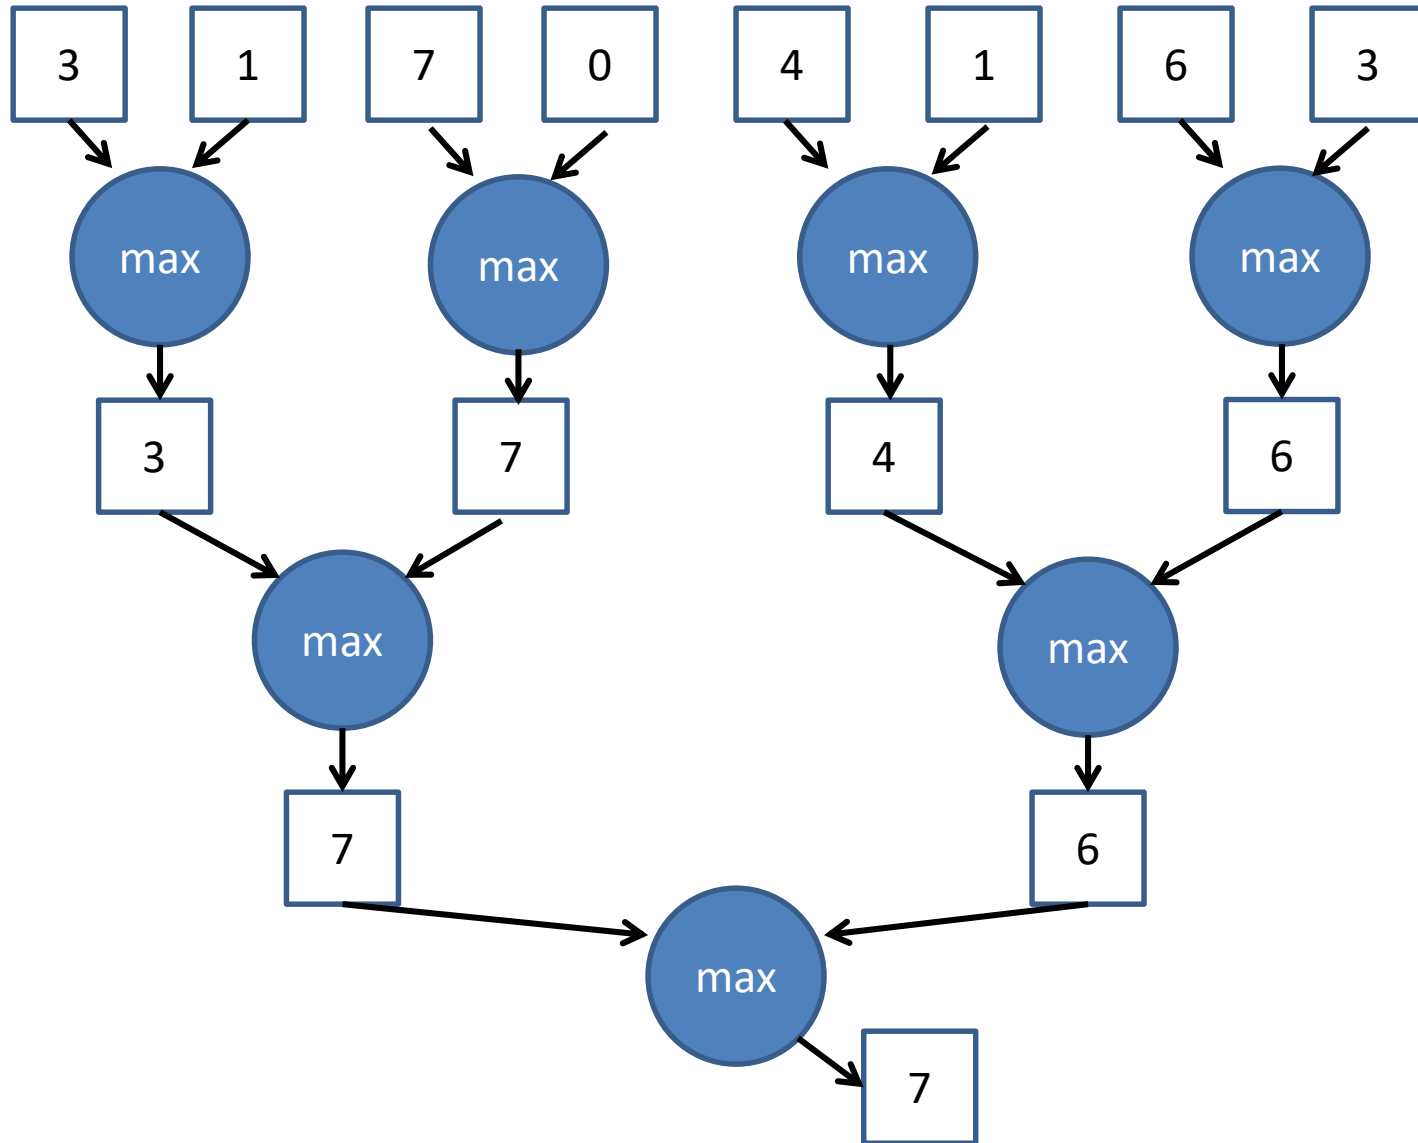
What is a reduction computation

- Summarize a set of input values into one value using a “reduction operation”
 - Max
 - Min
 - Sum
 - Product
 - Often with user defined reduction operation function as long as the operation
 - Is associative and commutative
 - Has a well-defined identity value (e.g., 0 for sum)

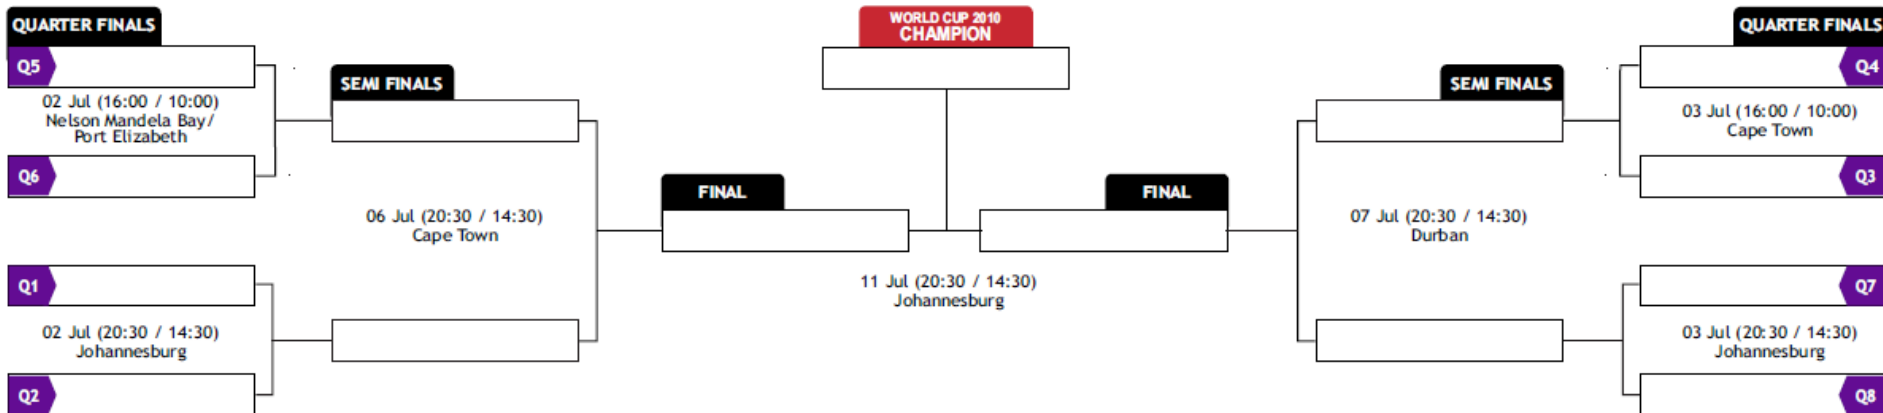
An efficient sequential reduction algorithm performs N operations in $O(N)$

- Initialize the result as an identity value for the reduction operation
 - Smallest possible value for max reduction
 - Largest possible value for min reduction
 - 0 for sum reduction
 - 1 for product reduction
- Scan through the input and perform the reduction operation between the result value and the current input value

A parallel reduction tree algorithm performs $N-1$
Operations in $\log(N)$ steps



A tournament is a reduction tree with "max" operation



A more artful rendition of the reduction tree.


A Quick Analysis

- For N input values, the reduction tree performs
 - $(1/2)N + (1/4)N + (1/8)N + \dots (1/N) = (1 - (1/N))N = N - 1$ operations
 - In $\log(N)$ steps - 1,000,000 input values take 20 steps
 - Assuming that we have enough execution resources
 - Average Parallelism $(N-1)/\log(N)$
 - For $N = 1,000,000$, average parallelism is 50,000
 - However, peak resource requirement is 500,000!
- This is a **work-efficient parallel algorithm**
 - The amount of work done is comparable to sequential
 - Many parallel algorithms are not work efficient

A Sum Reduction Example

- Parallel implementation:
 - Recursively halve the # of threads, add two values per thread in each step
 - Takes $\log(n)$ steps for n elements, requires $n/2$ threads
- Assume an in-place reduction using shared memory
 - The original vector is in device global memory
 - The shared memory is used to hold a partial sum vector
 - Each step brings the partial sum vector closer to the sum
 - The final sum will be in element 0
 - Reduces global memory traffic due to partial sum values

Vector Reduction with Branch Divergence



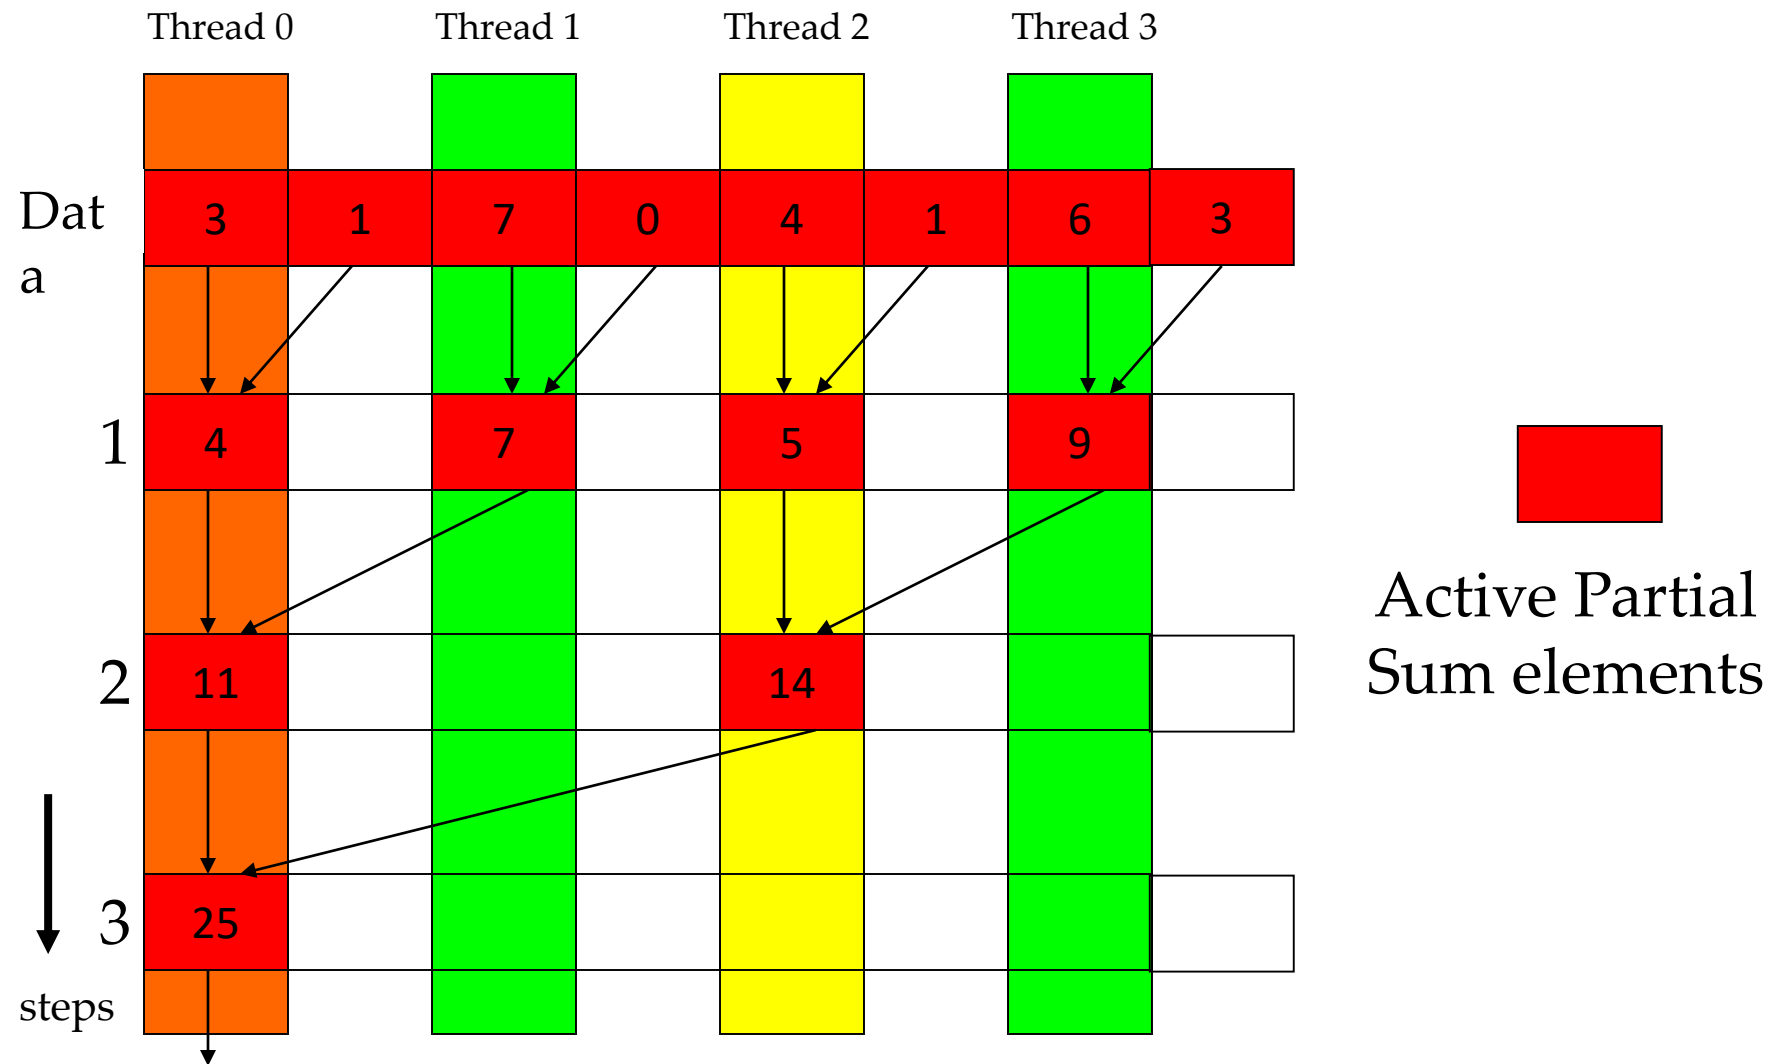
Simple Thread Index to Data Mapping

- Each thread is responsible of an even-index location of the partial sum vector
 - locations: 0, 2, 4, 6, ... hold sum of $0+1$, $2+3$, $4+5$, ...
- After each step, half of the threads are no longer needed
- In each step, one of the inputs comes from an increasing distance away

Optimizing Reduction Trees

- Performance factors of a reduction kernel
 - Memory coalescing
 - Control divergence
 - Thread utilization

A Sum Example (review)



The Reduction Steps

```
for (unsigned int stride = 1;
     stride <= blockDim.x; stride *= 2)
{
    __syncthreads();
    if (t % stride == 0) // t is thread ID
        partialSum[2*t] +=
        partialSum[2*t+stride];
}
```

Why do we need syncthreads()?

Barrier Synchronization

- `__syncthreads()` are needed to ensure that all elements of each version of partial sums have been generated before we proceed to the next step
- Why not another `__syncthread()` at the end of the reduction loop?

Back to the Global Picture

- At the end of the kernel execution, thread 0 in each block writes the sum of the block (stored in `partialSum[0]`) into a vector indexed by the value of `blockIdx.x`
- There can be a large number of such sums if the original vector is very large
 - The host code may iterate and launch another kernel
- If there are only a small number of sums, the host can simply transfer the data back and add them together.

Some Observations

- In each iteration, two control flow paths will be sequentially traversed for each warp
 - Threads that perform addition and threads that do not
 - Threads that do not perform addition still consume execution resources
- No more than half of threads will be executing after the first step
 - All odd-index threads are disabled after first step
 - After the 5th step, entire warps in each block will fail the if-condition, **poor resource utilization but no divergence**.
 - This can go on for a while, up to 5 more steps ($1024/32=16=2^5$), where each active warp only has one productive thread until all warps in a block retire

Thread Index Usage Matters

- In some algorithms, one can shift the index usage to improve the divergence behavior
 - Commutative and associative operators
 - At the end, the performance of many CUDA kernels depends on clever indexing.
- Reduction satisfies this criterion.


A Better Strategy

- Always compact the partial sums into the first locations in the partialSum[] array
- Keep the active threads consecutive

An Example of 16 threads

Thread 0 Thread 1 Thread 2

Thread 14Thread 15



A Better Reduction Kernel

```
for (unsigned int stride =  
    blockDim.x;  
    stride >= 1;  stride /= 2)  
{  
    __syncthreads();  
    if (t < stride) // t is thread ID  
        partialSum[t] +=  
        partialSum[t+stride];  
}
```

A Quick Analysis

- For a 1024 thread block
 - No divergence in the first 5 steps
 - 1024, 512, 256, 128, 64, 32 consecutive threads are active in each step
 - The final 5 steps will still have divergence

Parallel Algorithm Overhead

```
__shared__ float partialSum[2*BLOCK_SIZE];
```

```
unsigned int t = threadIdx.x;
```

```
unsigned int start = 2*blockIdx.x*blockDim.x;
```

```
partialSum[t] = input[start + t];
```

```
partialSum[blockDim+t] = input[start+ blockDim.x+t];
```

```
for (unsigned int stride = blockDim.x;
```

```
    stride >= 1;  stride >>= 1)
```

```
{
```

```
    __syncthreads();
```

```
    if (t < stride)
```

```
        partialSum[t] += partialSum[t+stride];
```


```
}
```

Parallel Algorithm Overhead


```
__shared__ float partialSum[2*BLOCK_SIZE];

unsigned int t = threadIdx.x;
unsigned int start = 2*blockIdx.x*blockDim.x;
partialSum[t] = input[start + t];
partialSum[blockDim+t] = input[start+ blockDim.x+t];
for (unsigned int stride = blockDim.x;
    stride >= 1;  stride >>= 1)
{
    __syncthreads();
    if (t < stride)
        partialSum[t] += partialSum[t+stride];
}
```

Parallel Execution Overhead




Parallel Execution Overhead



Although the number of “operations” is N , each “operation involves much more complex address calculation and intermediate result manipulation.

If the parallel code is executed on a single-thread hardware, it would be significantly slower than the code based on the original sequential algorithm.



Parallel Scan (Prefix Sum)

What? Why?

- Frequently used for parallel work assignment and resource allocation
- A **key primitive** in many parallel algorithms to convert serial computation into parallel computation
 - Based on reduction tree and reverse reduction tree

(Inclusive) Scan (Prefix-Sum) Definition

Definition: *The scan operation takes a binary associative operator \oplus , and an array of n elements*

$$[x_0, x_1, \dots, x_{n-1}],$$

and returns the prefix-sum array

$$[x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-1})].$$

Example: If \oplus is addition, then the scan operation on the array

[3 1 7 0 4 1 6 3]

would return

[3 4 11 11 15 16 22 25]

A Inclusive Scan Application Example

- Assume that we have a 100-inch bread to feed 10
- We know how much each person wants in inches
 - [3 5 2 7 28 4 3 0 8 1]
- How do we cut the bread quickly?
- How much will be left?
- Method 1: cut the sections sequentially: 3 inches first, 5 inches second, 2 inches third, etc.
- Method 2: calculate prefix-sum array
 - [3, 8, 10, 17, 45, 49, 52, 52, 60, 61] (39 inches left)
 - You can make 10 cuts in parallel at the above 10 cut points

Typical Applications of Scan

- Scan is a simple and useful parallel building block
 - Convert recurrences from sequential :

```
for(j=1;j<n;j++) out[j] = out[j-1] + f(j);
```
 - into parallel:

```
forall(j) { temp[j] = f(j) };  
scan(out, temp);
```
- Useful for many parallel algorithms:
 - radix sort
 - quicksort
 - String comparison
 - Lexical analysis
 - Stream compaction
 - Polynomial evaluation
 - Solving recurrences
 - Tree operations
 - Histograms
 - ...

Other Applications

- Assigning camp slots
- Assigning farmer market space
- Allocating memory to parallel threads
- Allocating memory buffer to communication channels
- ...

An Inclusive Sequential Scan

Given a sequence $[x_0, x_1, x_2, \dots]$

Calculate output $[y_0, y_1, y_2, \dots]$

Such that

$$y_0 = x_0$$

$$y_1 = x_0 + x_1$$

$$y_2 = x_0 + x_1 + x_2$$

...

Using a recursive definition

$$y_i = y_{i-1} + x_i$$

A Sequential C Implementation

```
y[0] = x[0];
```

```
for (i = 1; i < Max_i; i++) y[i] = y[i-1] + x[i];
```

Computationally efficient:

N additions needed for N elements → $O(N)$

A Naïve Inclusive Parallel Scan

- Assign one thread to calculate each y element
- Have every thread to add up all x elements needed for the y element

$$y_0 = x_0$$

$$y_1 = x_0 + x_1$$

$$y_2 = x_0 + x_1 + x_2$$

"Parallel programming is easy as long as you do not care about performance."

Parallel Inclusive Scan using Reduction Trees

- Calculate each output element as the reduction of all previous elements
 - Some reduction partial sums will be shared among the calculation of output elements
 - Based on design by Peter Kogge and Harold Stone at IBM in the 1970s - **Kogge-Stone Trees**


A Slightly Better Parallel Inclusive Scan Algorithm

T	3	1	7	0	4	1	6	3
---	---	---	---	---	---	---	---	---

1. Load input from global memory into shared memory array T

Each thread loads one value from the input (global memory) array into shared memory array T.

A Kogge-Stone Parallel Scan Algorithm




1. (previous slide)
2. Iterate $\log(n)$ times, **stride from 1 to $\text{ceil}(n/2.0)$** . Threads from stride to $n-1$ are active: add pairs of elements that are *stride* elements apart.

Iteration #1
Stride = 1

- Active threads: *stride* to $n-1$ ($n-\text{stride}$ threads)
- Thread j adds elements j and $j-\text{stride}$ from T and writes result into shared memory buffer T
- Each iteration requires two synctreads
 - `synctreads();` // make sure that input is in place
 - `float temp = T[j] + T[j - stride];`
 - `synctreads();` // make sure that previous output has been consumed
 - `T[j] = temp;`

A Kogge-Stone Parallel Scan Algorithm




1. ...

2. Iterate $\log(n)$ times, stride from 1 to $\text{ceil}(n/2.0)$. Threads *stride* to $n-1$ active: add pairs of elements that are *stride* elements apart.

Iteration #2
Stride = 2

A Kogge-Stone Parallel Scan Algorithm



1. Load input from global memory to shared memory.
2. Iterate $\log(n)$ times, stride from 1 to $\text{ceil}(n/2.0)$. Threads *stride* to $n-1$ active: add pairs of elements that are *stride* elements apart.
3. Write output from shared memory to device memory

Iteration #3
Stride = 4

Enhancement: Double Buffering

- Use two copies of data T0 and T1
- Start by using T0 as input and T1 as output
- Switch input/output roles after each iteration
 - Iteration 0: T0 as input and T1 as output
 - Iteration 1: T1 as input and T0 as output
 - Iteration 2: T0 as input and T1 as output
- This is typically implemented with two pointers, source and destination that swap their contents from one iteration to the next
- This eliminates the need for the second syncthread


Work Efficiency Analysis

- A Kogge-Stone scan kernel executes $\log(n)$ parallel iterations
 - The steps do $(n-1), (n-2), (n-4), \dots, (n - n/2)$ add operations each
 - Total # of add operations: $n * \log(n) - (n-1) \rightarrow O(n * \log(n))$ work
- This scan algorithm is not very work efficient
 - Sequential scan algorithm does n adds
 - A factor of $\log(n)$ hurts: 20x for 1,000,000 elements!
 - Typically used within each block, where $n \leq 1,024$
- A parallel algorithm can be slow when execution resources are saturated due to low work efficiency

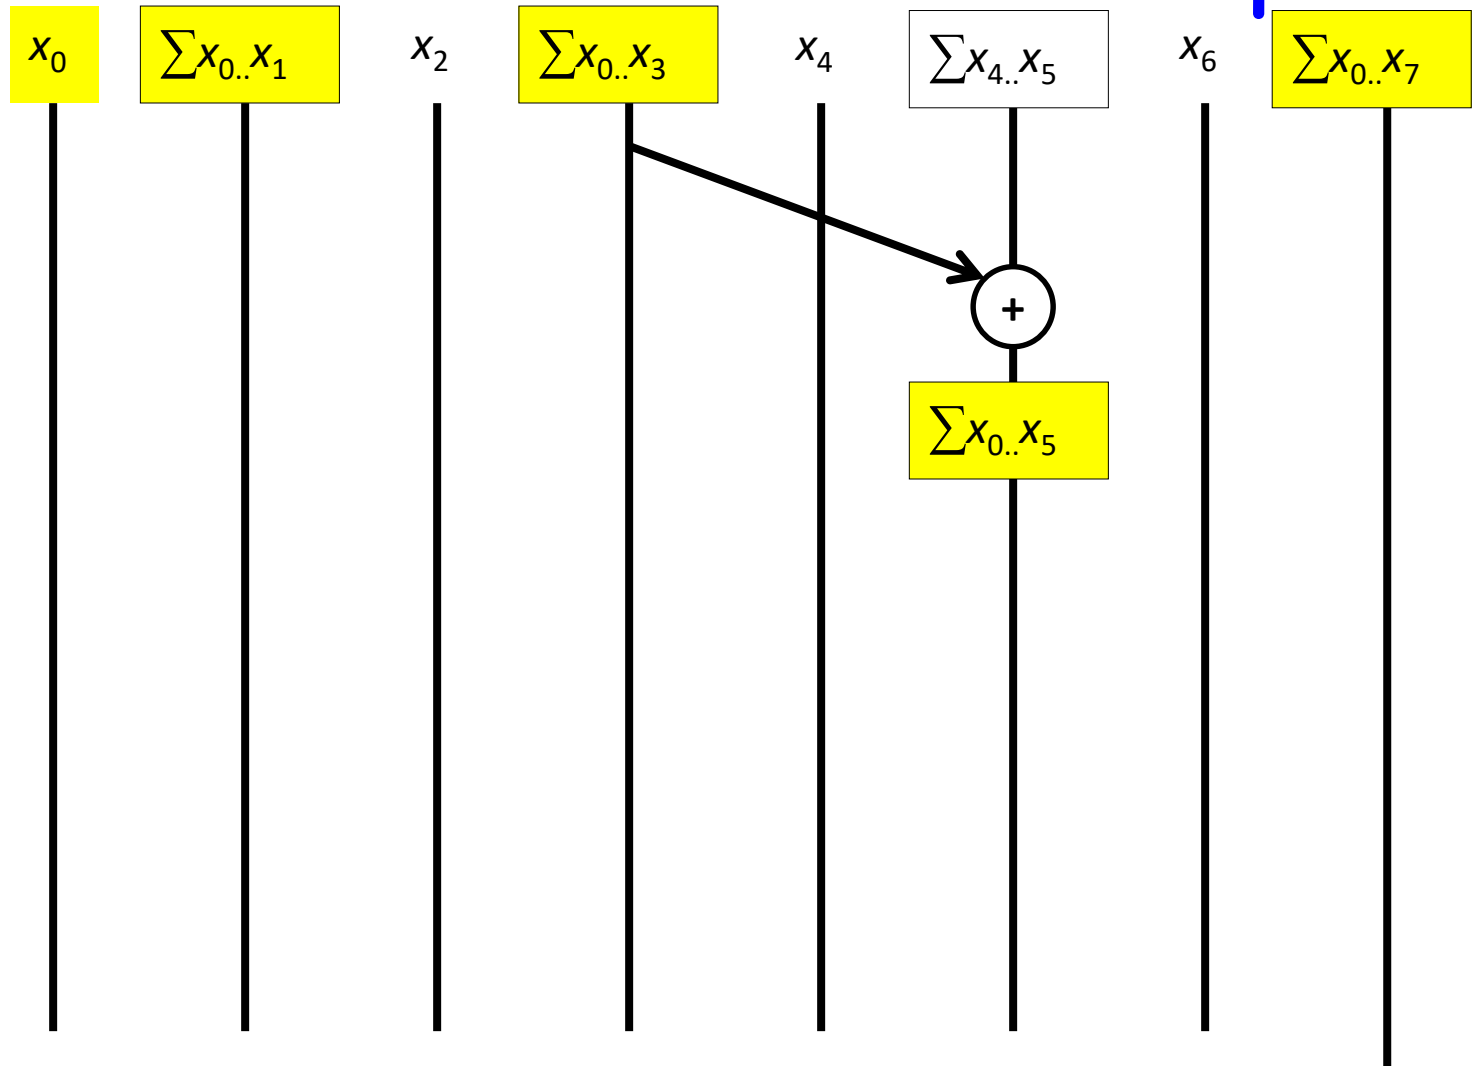
Improving Efficiency

- A common parallel algorithm pattern:
Balanced Trees
 - Build a balanced binary tree on the input data and sweep it to and from the root
 - Tree is not an actual data structure, but a concept to determine what each thread does at each step
- For scan:
 1. Traverse down from leaves to root building partial sums at internal nodes in the tree
 - Root holds sum of all leaves
 2. Traverse back up the tree building the scan from the partial sums

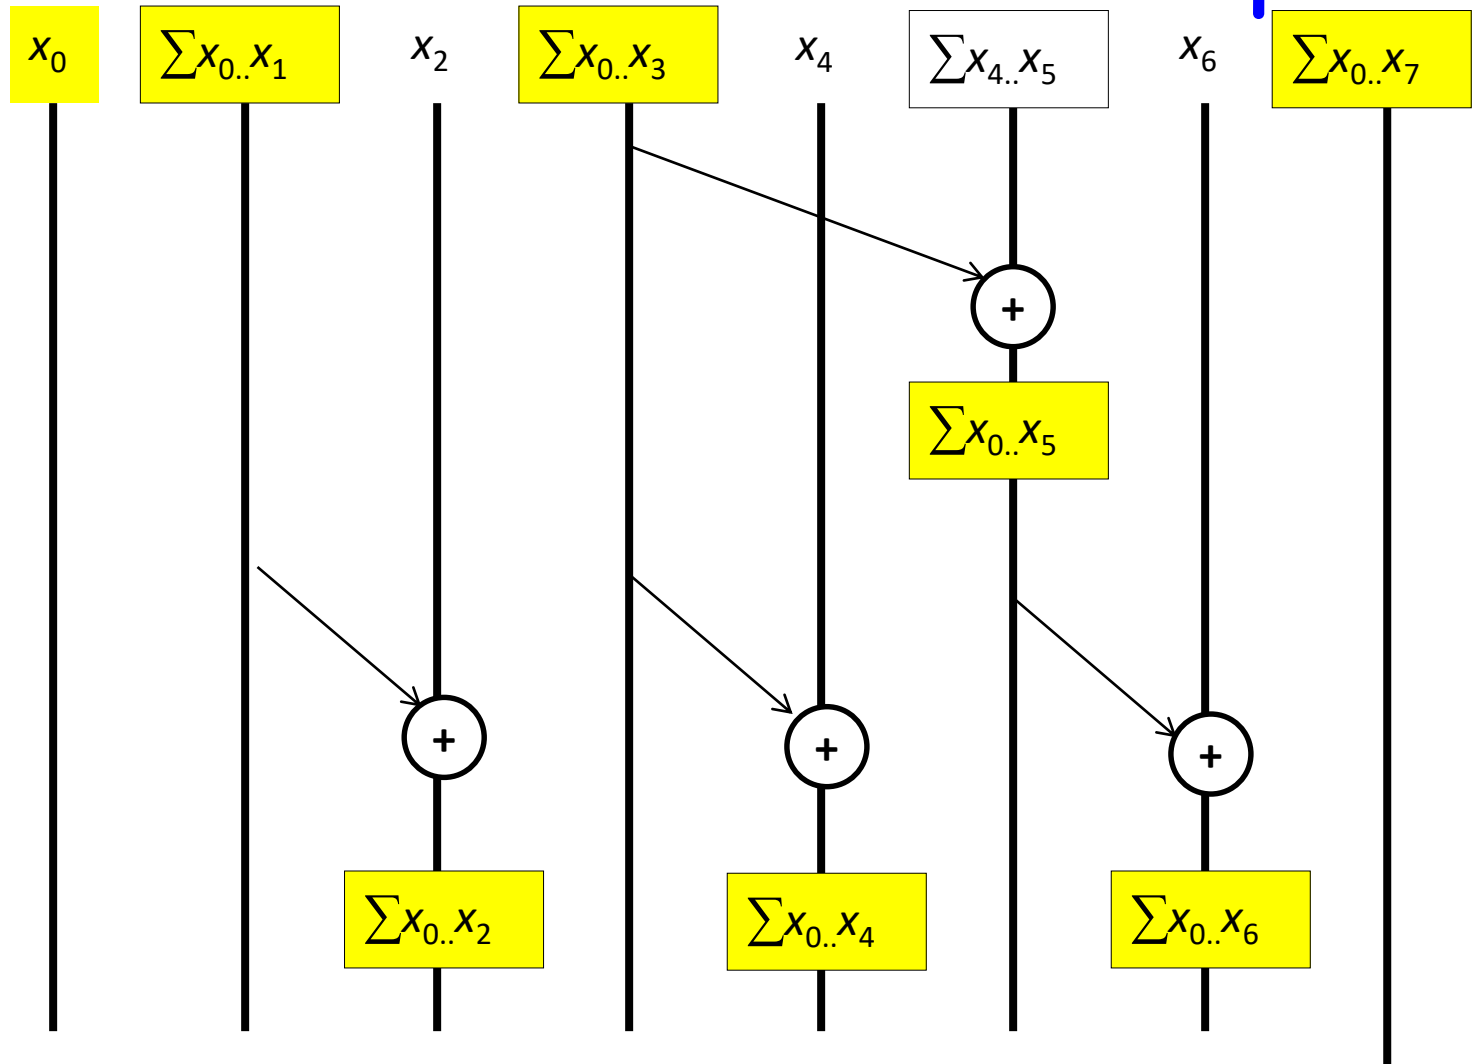
Brent-Kung Parallel Scan - Reduction Step




Inclusive Post Scan Step



Inclusive Post Scan Step



Putting it Together



Reduction Step Kernel Code

```
// float T[BLOCK_SIZE] is in shared memory
```

```
int stride = 1;
while(stride < BLOCK_SIZE)
{
    int index = (threadIdx.x+1)*stride*2 - 1;
    if(index < BLOCK_SIZE)
        T[index] += T[index-stride];
    stride = stride*2;

    __syncthreads();
}
```

Post Scan Step

```
int stride = BLOCK_SIZE/2;
while(stride > 0)
{
    int index = (threadIdx.x+1)*stride*2 - 1;
    if(index < BLOCK_SIZE)
    {
        T[index+stride] += T[index];
    }
    stride = stride / 2;
    __syncthreads();
}
```

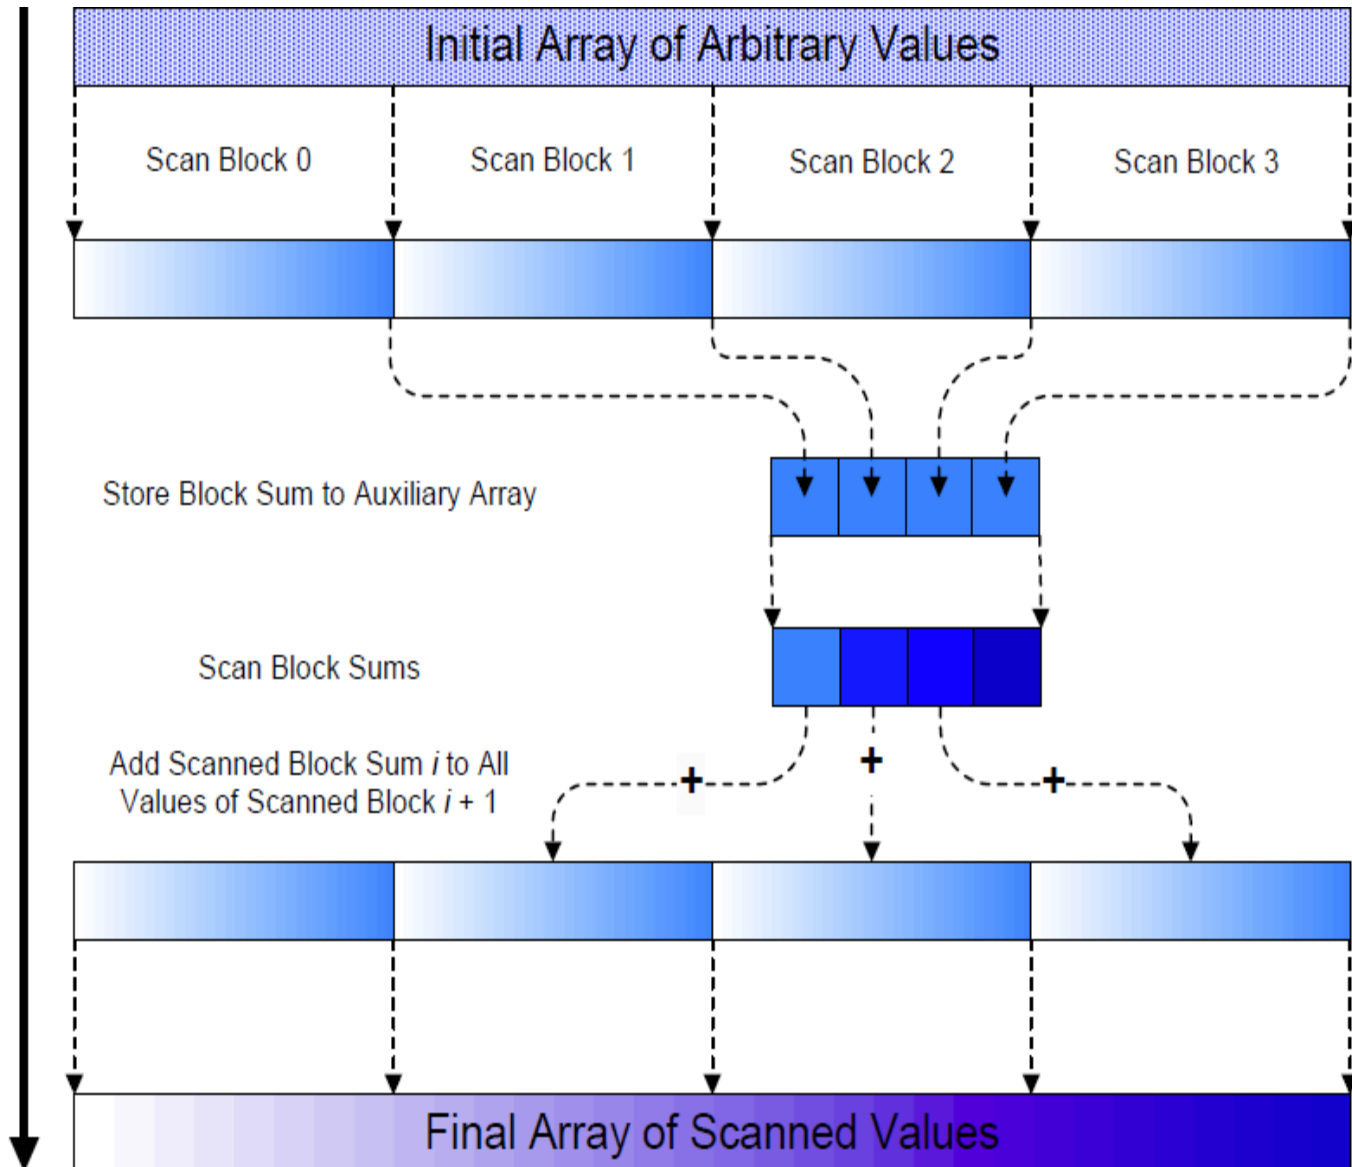
Work Analysis

- The parallel Inclusive Scan executes $2 * \log(n)$ parallel iterations
 - $\log(n)$ in reduction and $\log(n)$ in post scan
 - The iterations do $n/2, n/4, \dots, 1, 1, \dots, n/4, n/2$ adds
 - Total adds: $2 * (n-1) \rightarrow O(n)$ work
- The total number of adds is no more than twice of that done in the efficient sequential algorithm
 - The benefit of parallelism can easily overcome the 2X work when there is sufficient hardware

A couple of details

- Brent-Kung uses half the number of threads compared to Kogge-Stone
 - Each thread should load two elements into the shared memory
- Brent-Kung takes twice the number of steps compared to Kogge-Stone
 - Kogge-Stone is more popular for parallel scan with blocks in GPUs

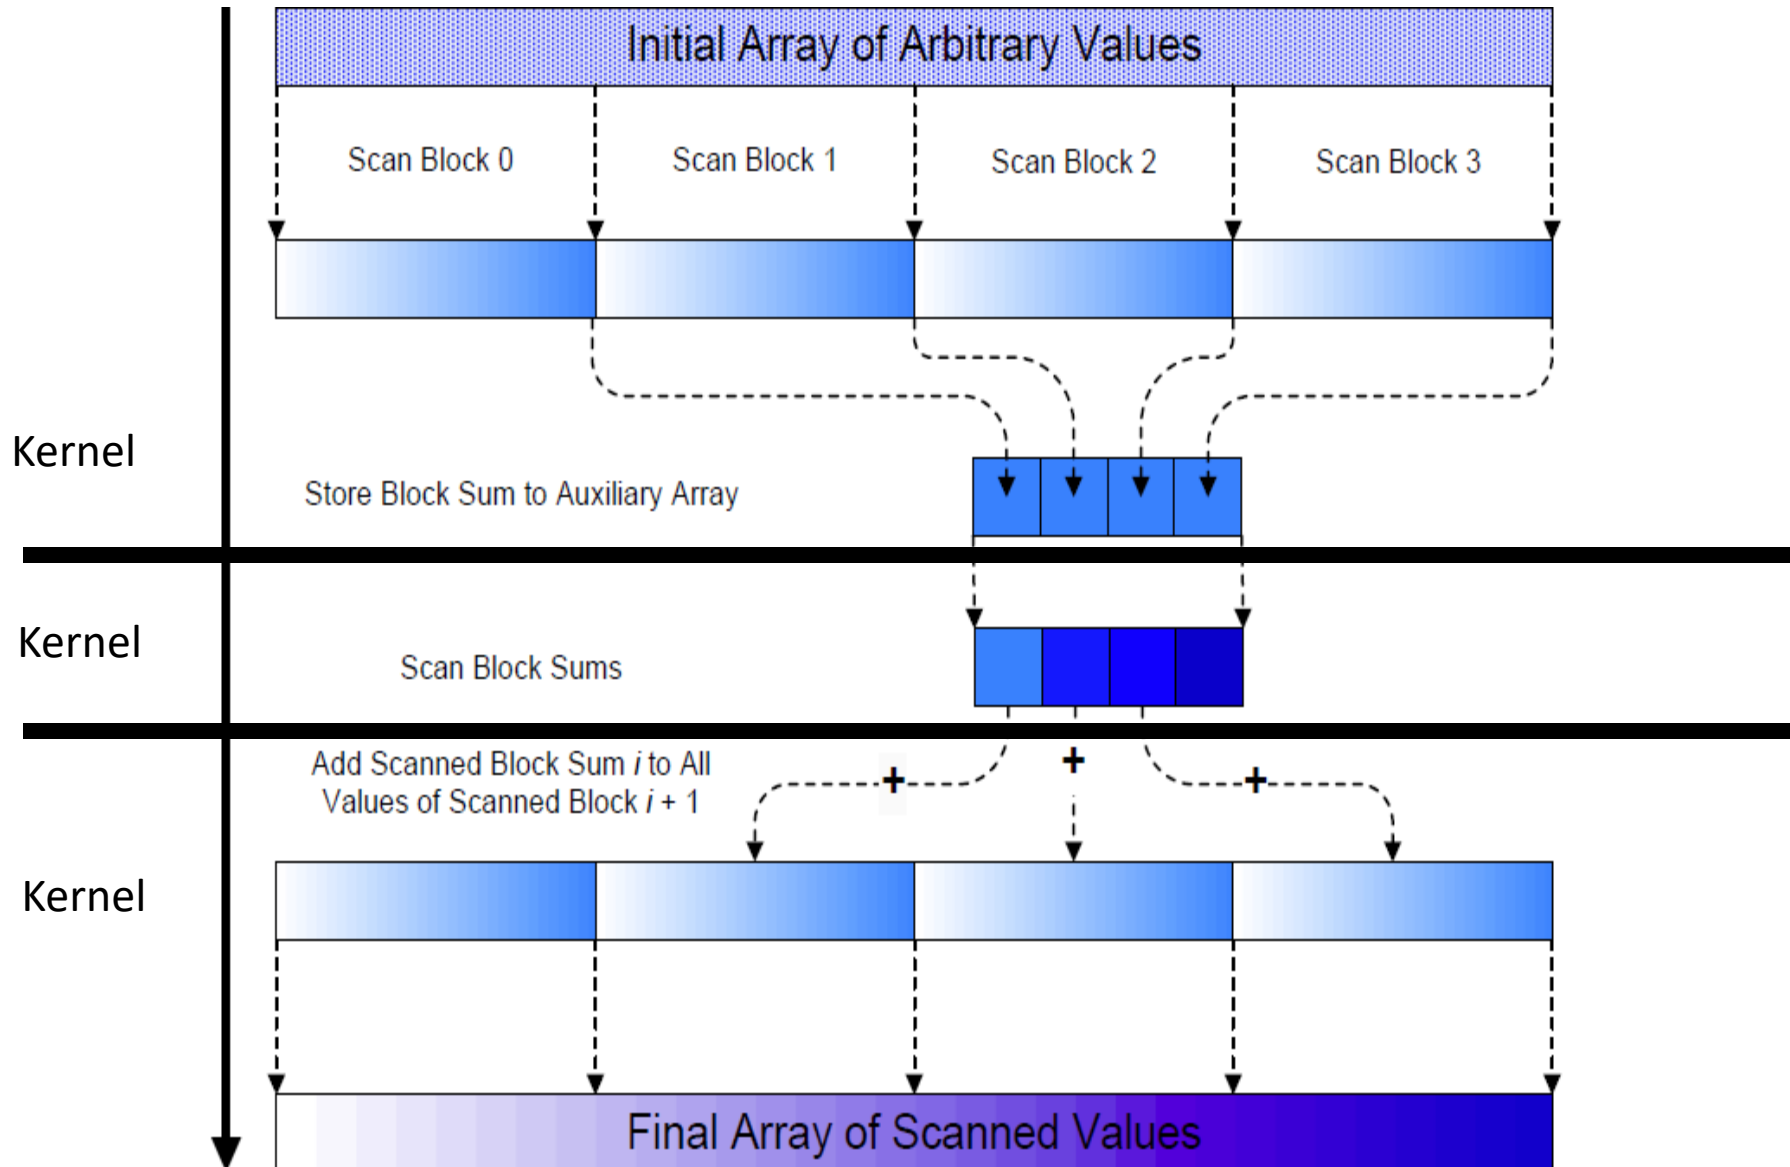
Overall Flow of Complete Scan A Hierarchical Approach



Using Global Memory Contents in CUDA

- Data in registers and shared memory of one thread block are not visible to other blocks
- To make data visible, the data has to be written into global memory
- However, any data written to the global memory are not visible until a memory fence. This is typically done by terminating the kernel execution
- Launch another kernel to continue the execution. The global memory writes done by the terminated kernels are visible to all thread blocks.

Overall Flow of Complete Scan A Hierarchical Approach



(Exclusive) Scan Definition

Definition: *The exclusive scan operation takes a binary associative operator \oplus , and an array of n elements*

$$[x_0, x_1, \dots, x_{n-1}]$$

and returns the array

$$[0, x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-2})].$$

Example: If \oplus is addition, then the exclusive scan operation

on $[3 \ 1 \ 7 \ 0 \ 4 \ 1 \ 6 \ 3]$

would return $[0 \ 3 \ 4 \ 11 \ 11 \ 15 \ 16 \ 22]$

Why Exclusive Scan

- To find the beginning address of allocated buffers
- Inclusive and Exclusive scans can be easily derived from each other; it is a matter of convenience

[3 1 7 0 4 1 6 3]

Exclusive [0 3 4 11 11 15 16 22]

Inclusive [3 4 11 11 15 16 22 25]

Conclusions

- We have reviewed several useful parallel patterns that you can use in your own GPU programming:
 - Convolution and tiled convolution
 - Reduction trees
 - Prefix scan (inclusive and exclusive)
- Parallel version must be work efficient
- Then we apply different GPU optimizations from our bag of tricks (coalescing, shared memory usage, ...).