

**CSCI-GA.3033-004**  
**Graphics Processing Units (GPUs): Architecture and Programming**  
**Homework 1**

(total: 25 points)

1. [11 points (0.5 per table entry)] The following are not exhaustive answers. You can come up with other, but correct answers too.

• Scenario	• Pros	• Cons
• Increasing number of SM	• More parallelism	• More pressure on L2 cache
• Increasing number of SPs (or cuda cores) per SM	• More opportunities for thread interactions • More parallelism • More resources = more blocks assigned to that SM	• Pressure on shared memory • Pressure on L1 cache
• Increasing memory bandwidth	• More efficient data supply to threads	• Wasted resource if not enough parallelism
• Increasing shared memory per SM	• More sharing among threads leading to less bandwidth requirement to global memory	• Wasted resources if not much data are shared
• Increasing L2 cache size	• L2 by definition coalesces memory accesses	• Many L2 cache misses consumes a lot of global memory bandwidth.
• Increasing number of SM • Increasing number of SPs (or cuda cores) per SM	• More parallelism	• More power consumption • More pressure on L2 caches
• Increasing number of SPs (or cuda cores) per SM • Increasing memory bandwidth	• More parallelism • More opportunities for memory coalescing	• More pressure on L1 cache • More pressure on shared memory • Wasted resources if not enough parallelism
• Increasing number of SM • Increasing shared memory per SM	• More parallelism	• More pressure on L2 cache
• Increasing number of SPs (or cuda cores) per SM • Increasing shared memory per SM	• More parallelism	• More pressure on L1 cache
• Increasing L2 cache size • Increasing shared memory per SM	• Less trips to global memory = less bandwidth requirement	• Wasted resources if not much data are shared • Bad L2 performance leads to more global memory access
• Increasing memory bandwidth • Increasing L2 cache size	• Can provide data faster to threads	• Wasted resources if not much data are shared

2. [2 points] Yes, because there is a big slowdown caused by moving the data from system memory to GPU memory. If we do not have enough data to ensure a lot of parallelism in the GPU, to overcome this performance loss, we may see *worse* performance for GPU over CPU-only version.

3. [8 points]

- a) Yes, each core can be assigned a subset of the numbers and check the existence of the number in them. This is done in parallel.
- b) Cannot be done on GPU because the operations are dependent on each other.
- c) The problem size is not big enough to ensure enough parallelism to overcome communication overhead.
- d) Yes, it is a highly parallel operation and we have enough independent computations. Moreover, the problem size is big enough to ensure enough parallelism.

4. [4 points]

- This will put a lot of pressure on that memory as it cannot serve all the cores in parallel. This memory will be a bottleneck that may affect scalability if we want to add more CPUs and/or GPUs.
- GPU needs a memory optimized for bandwidth while CPU needs memory optimized for latency.